

# Different Attention Mechanisms

Ritesh Kumar Maurya

2024-05-24

- In this blog, we will try to understand how different attention mechanisms work alongside vanilla attention from 2017 paper
- For the code part, you can checkout this [link](#)

## 0.1 Issue with RNNs

- Informaton Bottleneck:- If sentence is too much long, they needed to compress all the previous token information in just a single vector  $h$
- Sequential Nature:- All the tokens being processed one by one

## 0.2 Quotient rule

- Will be used while finding out the derivative of softmax
- $\left(\frac{u}{v}\right)' = \frac{u'v - uv'}{v^2}$ .

## 0.3 Transformers

- Given a sequence of length  $N$ , where each token is represented by  $D_{model}$  numbers, we compute three matrices
  - Query =  $XWq$  (what I'm looking for)
  - Key =  $XWk$  (what I have)
  - Value =  $XWv$  (the information I'll give if I'm relevant)
  - $Wq$  and  $Wk$  are of same shape and to maintain the symmetry  $Wv$  is also kept to be of same shape,  $Wq, Wk [D_{model}, D_k]$  and  $Wv [D_{model}, D_v]$
- Now to find out which key is more relevant to given queries we do matrix multiplication and the finally apply softmax to get probability
  - score =  $\text{Query} \text{Key}^\top$

- attn = softmax(score/sqrt(Dk))
  - Why scaling factor
    - \* this score dot product has variance of Dk, which causes instability in backpropagation
      - lets say we have two vectors q and k of dimension dk
      - $E[q_i] = E[k_i] = 0$
      - $\text{var}[q_i] = \text{var}[k_i] = 1$
      - then  $E[q \cdot k] = 0$  but  $\text{var}[q \cdot k] = dk$  [link](#)
    - \* the derivative of softmax contains some terms which if not normalized lead to gradient nearly 0 because of this larger variance
      - \*
- $$\text{Softmax}(z_i) = \frac{e^{z_i}}{\sum_j e^{z_j}}$$
- \*
- $$\frac{\partial \text{Softmax}(z)_i}{\partial z_j} = \text{Softmax}(z)_i (\delta_{ij} - \text{Softmax}(z)_j)$$
- \* If  $z_i$  is much larger than other elements,  $\text{Softmax}(z)_i \approx 1$ .
  - \* If  $z_i$  is much smaller,  $\text{Softmax}(z)_i \approx 0$ .
  - \* So in case of jacobian matrix, considering a simple case where query 0 attends all the keys
    - Diagonal:  $s_1(1 - s_1) \approx 1(0) = 0$ .
    - Off-Diagonal:  $-s_1s_2 \approx -1(0) = 0$ .
  - \* to enforce unit variance, we apply normalization by dividing the score by standard deviation i.e.  $\text{sqrt}(Dk)$

## 0.4 Single Head Vanilla Attention PyTorch Code

```

import torch
import torch.nn as nn
import torch.nn.functional as F
import time

device = torch.device("cuda" if torch.cuda.is_available() else "cpu")

class SingleHeadVanillaAttention(nn.Module):
    def __init__(self, in_features, out_features):
        super().__init__()
        self.in_features = in_features
        self.out_features = out_features
        self.Wq = nn.Linear(in_features = self.in_features, out_features = self.out_features)

```

```

        self.Wk = nn.Linear(in_features = self.in_features, out_features = self.out_features)
        self.Wv = nn.Linear(in_features = self.in_features, out_features = self.out_features)

    def forward(self, x):
        query = self.Wq(x) # [B, N, Dmodel] * [Dmodel, Dk] -> [B, N, Dk] TC=O(B*N*Dmodel*Dk)
        key = self.Wk(x) # [B, N, Dmodel] * [Dmodel, Dk] -> [B, N, Dk] TC=O(B*N*Dmodel*Dk)
        value = self.Wv(x) # [B, N, Dmodel] * [Dmodel, Dk] -> [B, N, Dk] TC=O(B*N*Dmodel*Dk)

        scores = torch.matmul(query, key.transpose(-1, -2))/(self.out_features**0.5) # [B, N, N]

        attn = F.softmax(scores, dim=-1) # [B, N, N]

        final = torch.matmul(attn, value) # [B, N, N] * [B, N, Dk] -> [B, N, Dk] TC=O(B*N^2*Dk)
        return final # [B, N, Dk]

Dmodel = 4096
Dk = 4096

model = SingleHeadVanillaAttention(in_features = Dmodel, out_features = Dk).requires_grad_(False)

n_tokens = 2048
inp = torch.randn(1, n_tokens, Dmodel, device=device)

torch.cuda.synchronize()
start = time.time()
k=50
for _ in range(k):
    with torch.no_grad():
        out = model(inp)
torch.cuda.synchronize()
print(f"Time taken: {(time.time()-start)/k:.3f}")

```

Time taken: 0.285

## 0.5 Linear Attention

- Instead of using softmax, it tries to make it more general where this softmax function can be replaced by any other function
- Kernel is some function which takes in two input vectors (query and key) and gives a single scalar output i.e. similarity score.

- This softmax can also be replaced by some kernel which is some kind of similarity function (i.e. to have property of a similarity function).
- Kernel can be decomposed into linear functions (feature functions)
  - $K(a, b) = \phi(a)\phi(b)^T$
- Feature function takes a vector as input and projects it into new feature space
- The basic idea is instead of doing complicated non-linear function like softmax, can we just project a and b into higher dimensional space and just do the linear inner product
- The attention mechanism is defined as  $A_l(x) = V' = \text{softmax}\left(\frac{QK^T}{\sqrt{D}}\right)V$ .
- We can make the above equation more generalised by using a similarity function
- $V'_i = \frac{\sum_{j=1}^N \text{sim}(Q_i, K_j)V_j}{\sum_{j=1}^N \text{sim}(Q_i, K_j)}$
- where we can think of  $\text{sim}(q, k) = \exp\left(\frac{q^T k}{\sqrt{D}}\right)$
- only constraint on sim is to give non-negative outputs
- given a kernel with feature representation  $\phi(x)$ , we can rewrite the above equation
- $V'_i = \frac{\sum_{j=1}^N \phi(Q_i)^T \phi(K_j)V_j}{\sum_{j=1}^N \phi(Q_i)^T \phi(K_j)}$
- By making the use of associative property of matrix multiplication, we can take out Q from summation, because it is independent of j
- $V'_i = \frac{\phi(Q_i)^T \sum_{j=1}^N \phi(K_j)V_j^T}{\phi(Q_i)^T \sum_{j=1}^N \phi(K_j)} \quad (4)$
- And finally we can convert it into vectorized format
- $(\phi(Q)\phi(K)^T)V = \phi(Q)(\phi(K)^TV)$
- It is linear because we can compute numerator and denominator and use it for all the queries [The eqn before vectorized one]
- For softmax attention, complexity is  $O(N^2 * \max(D, M))$ , where D is dimensionality for query and key, M is for value
- For linear attention, complexity is  $O(NCM)$ , where C is feature map projection dim and M is for value
- Linearization of exact softmax attention is not feasible because of the exponential kernel, which is of infinite dimension and can't be stored on a computer (expansion of  $e^x$  is infinitely many terms)
- That's why it is better to use some kind of polynomial feature representation

- The author chose elu instead of relu to avoid setting gradients to 0 when input was negative